

Classification


Tiphaine Viard

What is classification?

Classification is the problem of identifying to which of a set of categories (sub-populations) a new observation belongs, on the basis of a training set of data containing observations (or instances) whose category membership is known.

Classification

Generalize known structures to apply to new data.



An e-mail program might attempt to classify an e-mail as “legitimate” or as “spam”.

Spam example

Data set that describes e-mail features for deciding if it is spam.

Example

Contains “Money”	Domain type	Has attach.	Time received	spam
yes	com	yes	night	yes
yes	edu	no	night	yes
no	com	yes	night	yes
no	edu	no	day	no
no	com	no	day	no
yes	cat	no	day	yes

Spam example

Data set that describes e-mail features for deciding if it is spam.

Example

Contains “Money”	Domain type	Has attach.	Time received	spam
yes	com	yes	night	yes
yes	edu	no	night	yes
no	com	yes	night	yes
no	edu	no	day	no
no	com	no	day	no
yes	cat	no	day	yes

Assume we have to classify the following new instance:

Contains “Money”	Domain type	Has attach.	Time received	spam
yes	edu	yes	day	?

Definition

Given a set of classes $C_1 \dots C_N$, a classifier algorithm builds a model that predicts for every unlabelled instance I the class C_i to which it belongs with accuracy.

Example

Spam filter

Example

Twitter Sentiment analysis: analyze tweets with positive or negative feelings

Example

Cat or Dog?

A note about classification

Finding a defining feature for classification can cause unjust consequences; **redlining** can be exacerbated by machine learning models.

Dutch welfare ordered to stop using SyRI (2020)

Find the balance between
technical and **social** aspects

Privacy

Discrimination

Opacity

...

Symbolic AI

Baked-in fairness

“White-box” models

...

Basic Classifiers

Training

Compute the majority class in the dataset

Prediction

Output the majority class


k -Nearest Neighbors (k -NN)

Training

Store all instance (+ eventual index)

Prediction

Find the k closest point in the input and output the majority over those k points.



k -Nearest Neighbors (k -NN)

Training

Store all instance (+ eventual index)

Prediction

Find the k closest point in the input and output the majority over those k points.

Closest according to what metric?

L_1 vs L_2 vs L_∞ vs COS

k -Nearest Neighbors (k -NN)

Training

Store all instance (+ eventual index)

Prediction

Find the k closest point in the input and output the majority over those k points.

Closest according to what metric?

L_1 vs L_2 vs L_∞ vs COS

Rule of thumb: $k = \sqrt{n}$

Problem. Finding distance from all points to all others is costly : $\mathcal{O}(n^2)$ time

We only care about close data points

Problem. Finding distance from all points to all others is costly : $\mathcal{O}(n^2)$ time

We only care about close data points

Solution : use **locality-sensitive hashing** (LSH)


- Close points in the same bucket with high probability,
- Distant points are in different buckets with high probability.

Random hyperplanes for LSH

```
1 Generate  $K$  hyperplanes  $h_1, \dots, h_K$ ;  
2  $\mathcal{H}_{i,k} \leftarrow -1, \forall i, \forall k$ ;  
3 for every datapoint  $x_i$  do  
4   if  $x_i \cdot h_k \geq 0$  then  
5      $\mathcal{H}_{i,k} \leftarrow 0$ ;  
6   end  
7   else  
8      $\mathcal{H}_{i,k} \leftarrow 1$ ;  
9   end  
10 end
```


Repeat multiple times \rightarrow more robust

LSH : an example



Hash	Data points
000	A, B
001	
010	C
011	
100	D
101	
110	
111	E, F, G

LSH : an example



Hash	Data points
000	A, B, C
001	
010	
011	
100	D
101	
110	
111	E, F, G

Formula

$$\frac{P(A) \times P(B|A)}{P(B)} = P(A|B)$$

Proof.

$$P(A \cap B) = P(A) \times P(B|A)$$

Proof.

$$P(A \cap B) = P(A) \times P(B|A)$$

$$P(A \cap B) = P(B) \times P(A|B)$$

Proof.

$$P(A \cap B) = P(A) \times P(B|A)$$

$$P(A \cap B) = P(B) \times P(A|B)$$

$$P(A) \times P(B|A) = P(B) \times P(A|B)$$

Bayes theorem

Proof.

$$P(A \cap B) = P(A) \times P(B|A)$$

$$P(A \cap B) = P(B) \times P(A|B)$$

$$P(A) \times P(B|A) = P(B) \times P(A|B)$$

$$\frac{P(A) \times P(B|A)}{P(B)} = P(A|B)$$



Bayes theorem

Formula

$$\frac{P(A) \times P(B|A)}{P(B)} = P(A|B)$$

Interpretation

$$\text{prior} \times \frac{\text{likelihood}}{\text{evidence}} = \text{posterior}$$

Grouping attributes

$$P(C_i) \times \frac{P(\bar{x}|C_i)}{P(\bar{x})} = P(C_i|\bar{x})$$

Multiple attributes

$$P(C_i) \times \frac{\prod_j P(x_j|C_i)}{P(\bar{x})} = P(C_i|\bar{x})$$

Grouping attributes

$$P(C_i) \times \frac{P(\bar{x}|C_i)}{P(\bar{x})} = P(C_i|\bar{x})$$

Multiple attributes

$$P(C_i) \times \frac{\prod_j P(x_j|C_i)}{P(\bar{x})} = P(C_i|\bar{x})$$

With independence hypothesis!

Grouping attributes

$$P(C_i) \times \frac{P(\bar{x}|C_i)}{P(\bar{x})} = P(C_i|\bar{x})$$

Multiple attributes

$$P(C_i) \times \frac{\prod_j P(x_j|C_i)}{P(\bar{x})} = P(C_i|\bar{x})$$

With independence hypothesis!

$P(\bar{x})$ does not change with the class

Tree Methods

Classification

Data set that describes e-mail features for deciding if it is spam.

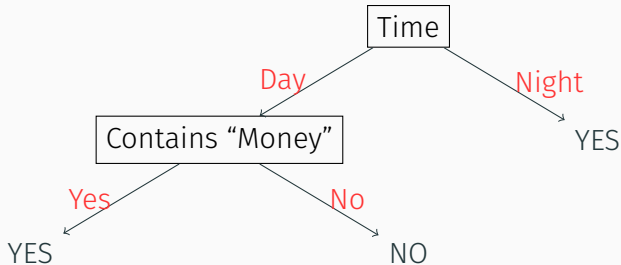
Example

Contains "Money"	Domain type	Has attach.	Time received	spam
yes	com	yes	night	yes
yes	edu	no	night	yes
no	com	yes	night	yes
no	edu	no	day	no
no	com	no	day	no
yes	cat	no	day	yes


Classification

Assume we have to classify the following new instance:

Contains "Money"	Domain type	Has attach.	Time received	spam
yes	edu	yes	day	?



Decision Trees



Recursive construction technique

- $A \leftarrow$ the *best* decision attribute for next *node*
- Assign A as decision attribute for *node*
- For each value of A , create new descendant of *node*
- Sort training examples to leaf nodes
- If training examples perfectly classified, Then STOP, Else iterate over new leaf nodes

- Interpretable: the DT is easy to understand
- Training is fast (greedy algorithm)
- Expressive: can approximate complex non-linear functions
- Complex function = large tree
- Large tree = more variance = overfitting

In practice, decision trees underperform compared with other methods

Bagging (Breiman, 1996)

Example

Dataset of 4 Instances : A, B, C, D

Classifier 1: B, A, C, B

Classifier 2: D, B, A, D

Classifier 3: B, A, C, B

Classifier 4: B, C, B, B

Classifier 5: D, C, A, C


Bagging:

1. Bootstrap: generate multiple samples of data, train decision tree
2. Aggregate: output the average output of all models

Bagging

Bagging builds a set of M base models, with a bootstrap sample created by drawing random samples with replacement.

- Highly expressive: each model can estimate complex functions/boundaries
- a low-variance method: averaging the prediction of all models reduces variance (if M large enough)



Evaluating bagging: out-of-bag error

Bagging implies multiple models, running on different, overlapping datasets

How can we evaluate its performance?

Evaluating bagging: out-of-bag error

Bagging implies **multiple models**, running on **different, overlapping** datasets

How can we evaluate its performance?

The **out-of-bag error**

1. for each data point x , average predicted output over models that do **not** contain x in bootstrap \Rightarrow **point-wise out-of-bag error**;
2. **Average** point-wise out-of-bag error over training set.

Problems with bagging

In practice, trees are **strongly correlated**.

- Suppose that x_j is a strong predictor. Then most models will split on x_j ;
- Then, each tree is essentially the same, and the averaged output is irrelevant.


For F identically, dependently distributed variables with pairwise correlation ρ and variance σ^2 , variance of mean:

$$\rho\sigma^2 + \frac{1-\rho}{F}\sigma^2$$

Variance reduction of bagging is **limited**.

Random Forests


- Bagging
- Random Trees: trees that in each node only uses a random subset of k of the attributes



⇒ one of the most popular methods in machine learning.

Feature importance with random forests

- Record prediction accuracy on the oob samples for each tree;
- Randomly permute the data for column j in the oob samples;
- Record the accuracy again;
- Average the decrease in accuracy over all trees, and use as a measure of the importance of variable j .



Final thoughts on Random Forest

- Ensemble methods: **not easily interpretable**
- One of the best “off-the-shelf” methods, near to 0 tuning necessary
- **Averaging** and **randomization** offer fine control of the bias-variance trade-off
- Reasonably efficient: $\Omega(mk\hat{n}\log^2\hat{n})$, with $\hat{n} \approx 0.63n$
- use **n_jobs** to make parallel
- **sklearn** implementation (Python + Cython) is by far the fastest available

Gradient Boosting


Train each of the M trees on the error of the precedent trees

$$\phi(x) = \sum_{m=0}^M \phi_m(x)$$

with each step being staged:

$$\phi_m(x) = \phi_{m-1}(x) + \hat{\phi}_m(x)$$

and $\hat{\phi}_m$ is a tree that approximates the gradient step.



Last thoughts on gradient boosting

- Usually more accurate than Random Forests
- Can adapt to most loss functions
- Under- and overfitting are addressed through regularization (learning rate, subsampling, term in loss function...)
- Tuning is harder than for random forests
- Slow to train (no parallelism!), fast predictions
- Still, blazing-fast implementation (Python and C++):
XGBoost
- Easy to get *feature importance*

Gradient-based Methods

Logistic Regression

Training

Learn an hyperplan \mathcal{P} separating well the two classes.

Prediction

What side of the hyperplan \mathcal{P} is the point?



Based on the gradient of the logit function.


Logistic Regression

Training

Learn an hyperplan \mathcal{P} separating well the two classes.

Prediction

What side of the hyperplan \mathcal{P} is the point?



Based on the gradient of the logit function.


Logistic Regression

Training

Learn an hyperplan \mathcal{P} separating well the two classes.

Prediction


What side of the hyperplan \mathcal{P} is the point?



Based on the gradient of the logit function.

Support-Vector Machines (SVM)


Similarly to logistic regression, we want to find the **best** separating hyperplane



SVM: under the hood

Suppose a linear separation exists

$$w^T x + b = 1, w^T x + b = -1$$




x_i near the margin determine the solution \rightarrow support vectors

SVM: under the hood

Suppose a linear separation exists

$$w^T x_i + b \geq 1, \text{ if } y_i = 1, w^T x_i + b \leq -1, \text{ if } y_i = -1$$




x_i near the margin determine the solution \rightarrow support vectors

SVM: under the hood

Suppose a linear separation exists

$$w^T x_i + b \geq 1, \text{ if } y_i = 1, \quad w^T x_i + b \leq -1, \text{ if } y_i = -1$$

$$y_i(w^T x + b) \geq 1$$




x_i near the margin determine the solution → support vectors

SVM: under the hood

Suppose a linear separation exists

minimize $\|w\|$ subject to $y_i(w^T x_i + b) \geq 1$, for $i = 1..n$



x_i near the margin determine the solution \rightarrow support vectors

minimize $\|w\|$ subject to $y_i(w^T x_i + b) \geq 1$, for $i = 1..n$

$\|w\|$ is the L_1 norm

The gradient of $\|w\|$ is $\frac{w}{\|w\|}$

Let's switch instead to L_2 norm

The gradient of $\|w\|^2$ is $2w$

minimize $\frac{1}{2}\|w\|^2$ subject to $y_i(w^T x_i + b) \geq 1$, for $i = 1..n$

This is a quadratic optimization problem.

- What about non linearly separable data?

Add misclassification term (number or distance)

- Use **hinge function** as regularization:

$$\min ||w||^2 + \lambda \left[\frac{1}{n} \sum_i^n \max(0, 1 - y_i(w^T x - b)) \right]$$

- λ trade-off between increasing margin and ensuring x_i outside of margin
- Only adapted to binary classification
- But you can do "one-versus-all" classification